

Composable Sparsity Gating for Completion-Aware Clock-Gated MAC Arrays

Kulamani Rout

Abstract—Operand zero-skipping and completion-aware clock gating are both established low-power techniques for multiply-accumulate (MAC) arrays, and are routinely described as complementary. This brief shows that they are not directly composable. Wherever a single clock enable is shared between the operation arriving this cycle and operations already in flight—the normal case for stage-level gating—the widely used sparsity enable $EN = valid \wedge (A \neq 0) \wedge (B \neq 0)$ silently applies one operation’s condition to its predecessors. With completion-aware (drain-counter) gating the consequence is deterministic rather than intermittent, because a correctly sized drain window has zero slack by construction and the stranded result is never produced. On an eight-lane cascaded MAC array this drops 63 % of all results. We give a structural fix—splitting the clock enable into a never-masked control enable and a sparsity-gated datapath enable, with the skip decision carried per item—and prove it cycle-for-cycle equivalent to the ungated reference over randomised sparse traffic, using the broken design as a negative control. We further show that register-enable activity counts, the metric usually quoted for such schemes, materially overstate the saving, and that an asynchronous datapath reset silently prevents the DSP48 register packing on which the real saving depends.

Index Terms—Clock gating, sparsity, multiply-accumulate, low-power design, FPGA, design verification.

I. INTRODUCTION

DYNAMIC power in a MAC array is dominated by the multiplier array and its pipeline registers. Two mitigations are standard. *Clock gating* suppresses register activity while a stage is idle [1], [2]. *Zero-skipping* exploits the high operand sparsity of quantised CNN inference: because $a = 0 \vee b = 0 \Rightarrow ab = 0$ exactly, the multiply may be skipped without any loss of accuracy. Zero-skip has been used at architectural scale in Eyeriss [3], Cnvlutin [4], SCNN [5] and EIE [6], is claimed in industrial patents [8], [9], and appears at RTL as the enable

$$EN_{MAC} = valid \wedge (A \neq 0) \wedge (B \neq 0) \quad (1)$$

used, for example, in a recent FPGA power-optimisation study [7].

The two techniques are habitually presented as orthogonal: one removes idle activity, the other removes ineffectual activity. This brief shows that the orthogonality fails as soon as the clock gate is *completion-aware*, that is, when the enable is held asserted by a drain counter until every in-flight pipeline slot has retired. The contributions are:

- 1) A demonstration that (1) is unsafe in a drain-gated pipeline, with the failure mechanism identified as the

zero-slack property of a correctly sized drain window (Section III).

- 2) A split-enable construction that restores composability at negligible cost, together with a lossless propagation of sparsity information into the downstream adder tree (Section IV).
- 3) An equivalence-based verification method that uses the broken design as a negative control, so that the passing result is shown to be non-vacuous (Section V).
- 4) Post-implementation evidence on both an FPGA and an open standard-cell flow that register-enable activity counts, the metric usually quoted for such schemes, can overstate the achievable saving by the entire multiplier contribution, and that the correction is affordable only if the datapath reset style permits register packing (Section VI).

II. COMPLETION-AWARE GATING BASELINE

The reference design (Fig. 1(a)) is an 8-lane array of two-stage (multiply, accumulate) 16-bit MACs feeding a three-level registered adder tree. Each stage owns an independent gating cell parameterised by its own pipeline depth D :

$$ce = pe \vee oa \vee trig, \quad (2)$$

where *trig* marks the cycle a new item enters, *oa* is a registered “operation active” flag, and *pe* is an external override. On *trig* a counter loads D ; otherwise it decrements, and *oa* falls when it reaches zero. The disjunction with *trig* in (2) is required because *oa* is registered and would otherwise drop the first cycle of every burst.

The array does not use a central controller. Stage 2 is triggered by the disjunction of the lanes’ completion outputs, so each stage independently wakes its successor—the arrangement is a cascade of autonomous gates rather than a scheduled pipeline. Related distributed pipeline gating is claimed in [10], which propagates stall and valid signals rather than maintaining per-stage drain counters.

III. THE COMPOSABILITY HAZARD

A. Drain windows have zero slack

A drain counter is correct when D equals the stage latency: a smaller D strands in-flight data, a larger one wastes enable cycles. Sizing it correctly therefore makes the enable window exactly as long as the pipeline needs and not one cycle longer. Measurement on the reference design confirms this. For a single item entering at cycle T , the MAC’s *ce* is asserted only during $\{T, T+1, T+2\}$, and *completion_out* fires in $T+2$, the *last* asserted cycle. The adder tree behaves identically over four cycles. There is no spare enable cycle anywhere.

The author is with [AFFILIATION TO BE SUPPLIED] (e-mail: routsarojini901@gmail.com).

Design sources, testbenches and the complete Vivado build/power scripts needed to reproduce every number in this brief are released with the paper.




Fig. 1. Architecture. (a) The array and the adder tree each own an independent completion-aware gate sized to their own pipeline depth D ; there is no central controller, and the update bus (blue) is the only wire the sparsity extension adds between stages. (b) Within a lane the clock enable is split. ce_{ctrl} is triggered by the raw *valid* and is never masked, so the drain counter is structurally blind to sparsity; ce_{mult} and ce_{acc} additionally carry the skip term and gate only the two wide register banks, which hold 64 of the lane's ≈ 70 flip-flops. The skip decision is registered as $skip_p$ alongside $valid_p$, so the accumulate stage consults the tag belonging to *its own* item rather than the current input bus.

B. Failure mechanism

Consider a burst in which a zero-valued operand arrives at $T+1$, one cycle after a non-zero item. Under (1) the shared enable is deasserted at $T+1$. That enable, however, is not private to the arriving item: it also clocks the accumulate stage of the item that entered at T and is mid-drain. Suppressing it defers that accumulation; because the drain counter continues to run regardless, the enable window closes before the deferred accumulation can occur, and the result is lost permanently. The three canonical cases are zero mid-burst, zero at burst start, and zero at burst end; only the last is fatal, and only because of the zero-slack property.

The hazard is more general than the drain counter that exposes it. It arises whenever a *single* clock enable is shared between the item arriving in the current cycle and items already in flight in the same stage—which is the normal situation for stage-level gating, however the enable is generated. Deriving that enable from a per-item condition such as (1) then silently applies one item's condition to its predecessors. What a correctly sized drain counter changes is only that it removes the accidental slack which might otherwise let the stranded operation catch up on a later idle cycle, and so converts an intermittent, data-dependent fault into a deterministic one. Schemes that hold the enable using a valid shift register, or a *busy/stall* signal, share the same exposure.

Fig. 2 contrasts the two cases. The effect is not marginal. Implementing (1) verbatim on the reference cascade and comparing against the ungated design under identical stimulus, the naive design produces **77 of 208** expected result pulses—**63 % of results silently dropped**—and differs from the reference on **278 of 284** compared cycles. It does not compute wrong sums so much as stop computing.




Fig. 2. The composability hazard. A zero operand at $T+2$ lands inside the drain window of the item that entered at $T+1$. (a) Masking the shared enable strands that item; because the drain counter runs regardless, the window closes and the result is never produced. (b) Splitting the enable leaves the drain path untouched and gates only the datapath, so the result is retained.

IV. SPLIT-ENABLE ZERO-SKIP

A. Construction

The fix is to stop treating the clock enable as a single signal (Fig. 1(b)). We derive

$$skip = valid \wedge (\overline{a} \vee \overline{b}) \quad (3)$$

$$ce_{ctrl} = f_{drain}(trig = valid) \quad (4)$$

$$ce_{mult} = ce_{ctrl} \wedge valid \wedge \overline{skip} \quad (5)$$

$$ce_{acc} = ce_{ctrl} \wedge valid_p \wedge \overline{skip_p} \quad (6)$$

Two properties matter. First, the drain counter is triggered by the raw *valid* and is therefore structurally blind to sparsity: no operand pattern can shorten, extend or corrupt it. Second, the skip decision is registered alongside the item as $skip_p$, exactly parallel to the validity tag $valid_p$, so the accumulate stage consults the tag belonging to *its own* item rather than whatever is on the input bus in the current cycle. This is what makes the zero-at-burst-end case safe: the arriving zero cannot suppress a predecessor's accumulation because that predecessor carries $skip_p = 0$.

TABLE I
EQUIVALENCE AGAINST THE UNGATED REFERENCE

Design	result pulses	diverging cycles	verdict
Reference (ungated)	208	—	—
Proposed (split enable)	208	0 / 284	pass
Naive, Eq. (1)	77	278 / 284	fail

Only the two wide register banks are sparsity-gated. The three control flip-flops per lane remain ungated; gating them would require perturbing completion timing, which the zero-slack window forbids, and they represent under 5% of the lane’s registers.

B. Hierarchical propagation

Per-lane skipping is local. The cascade admits a stronger statement: if a lane’s accumulator did not change, every partial sum in the tree that depends only on unchanged accumulators is also unchanged and its register may be held. Each lane exports *upd* (“my accumulator moved”), registered in the same cycle as its completion output, from which

$$en_j^{(1)} = upd_{2j} \vee upd_{2j+1}, \quad en_k^{(l+1)} = chg_{2k}^{(l)} \vee chg_{2k+1}^{(l)} \quad (7)$$

with $chg^{(l)}$ the registered form of $en^{(l)}$, so that the change flags advance up the tree in the same pipeline phase as the data they describe. Holding a register whose inputs provably did not change is bit-exact. Note the exploited property is *unchangedness*, not zeroness: the tree sums accumulators, which are persistent state and rarely zero. This distinguishes the mechanism from operand-zero gating of adder trees [9]. Under unstructured sparsity s a level-1 adder idles only when both feeding lanes idle, i.e. with probability s^2 ; the measured level-1 and level-2 savings at $s = 0.48$ are 22.7% and 4.9% against $s^2 = 23\%$ and $s^4 = 5.3\%$ predicted. The tree contribution is therefore genuinely second order and is reported as such.

V. VERIFICATION

Because zero-skip is lossless by construction, the gated and ungated designs must agree on *every* cycle, not merely at the end of a burst. Both are driven from identical stimulus wires and compared cycle-for-cycle on (*result*, *result_valid*); an independent behavioural model of the per-lane accumulators is checked in parallel, so a fault corrupting both designs identically would still be caught. Stimulus comprises eleven directed cases covering the three canonical zero positions, all-zero bursts, a twelve-cycle zero run with resumption, per-lane heterogeneous sparsity, and a 400-cycle randomised sparse stream.

A passing equivalence check proves nothing unless it can fail. We therefore retain the naive design of Section III as a *negative control* and run it in the same harness. Table I reports both. The reference emits 208 result pulses, establishing non-vacuity; the proposed design matches on every cycle; the naive design is caught.

TABLE II
OUT-OF-CONTEXT IMPLEMENTATION ON XC7A100TCSG324-1, 4.0 NS TARGET

Design	LUT	FF	CARRY	DSP regs [†] A/B/M/P	$F_{\max}$ (MHz)
<i>2-stage pipeline, asynchronous datapath reset</i>					
Baseline	789	546	127	0/0/0/8	269.8
+ split-enable zero-skip	890	568	127	0/0/0/8	272.7
<i>3-stage pipeline, asynchronous datapath reset</i>					
+ gated operand regs	906	840	127	0/0/0/8	202.6
<i>3-stage pipeline, no datapath reset</i>					
control, gating off	413	298	63	8/8/8/8	271.7
+ gated operand regs	530	328	63	8/8/8/8	283.0

[†]Number of the eight DSP48E1 slices whose A, B, M and P pipeline registers are used. All designs infer 8 DSP48E1 and no BUFGCE.

VI. IMPLEMENTATION RESULTS

All designs were synthesised and implemented out of context on xc7a100tcs324-1 with Vivado 2023.2 at a 4.0 ns target, budgeting 20% of the period to off-core delay on both inputs and outputs. Out-of-context implementation is required—the cascade has 311 ports against 210 available pins—and is also appropriate here because it excludes I/O buffer power from the comparison. Timing every port path uniformly matters: false-pathing the inputs would exclude the multiplier-input paths of the two-stage designs from analysis altogether and render the frequency column meaningless. Power is obtained from `report_power` driven by SAIF activity captured from simulation of the same LFSR stimulus, so every design sees bit-identical data at a given sparsity. Table II summarises area and timing.

A. The cost of composability

The final two rows of Table II differ only in whether sparsity gating is enabled, and so isolate its true cost: **+117 LUT (+28%) and +30 FF (+10%), with no frequency penalty** (283.0 against 271.7 MHz). Applied instead to the original two-stage pipeline (the two-stage rows) the split-enable construction costs +101 LUT (+12.8%) and +22 FF, again at no loss of $F_{\max}$. Composability is therefore cheap; what is expensive is getting the surrounding register style wrong, as the next two subsections show.

B. Register-enable counts overstate the saving

The two-stage split-enable design gates the *product* register, and activity counting there suggests a 64% reduction in load events at 30% sparsity. Post-implementation this is misleading. Vivado maps each multiply to a DSP48E1 but reports $AREG = BREG = MREG = 0$ with only the product register absorbed ($PREG = 8$). The multiplier array is consequently driven combinatorially from whatever lies outside the DSP—here, directly from the operand ports, which toggle every cycle irrespective of any enable. Gating the product register saves that register’s own load and nothing of the array behind it.

Quieting the array requires *gated operand* registers, so that a skipped item leaves the DSP inputs unchanged. Measured

at 29 % sparsity, transitions at the multiplier inputs fall from 2661 to 1009, a 62 % reduction, whereas the product-register-gated design scores *zero* on the same metric. The distinction is invisible to register-enable activity counting. We suggest that where such counts are used to justify a gating scheme, the register being gated should be stated explicitly, since gating a pipeline register downstream of a combinational block does not quiet that block.

C. Asynchronous reset blocks the fix

Adding the operand registers in the design's existing house style—with asynchronous reset—leaves them in fabric (third block of Table II). The critical path becomes a fabric flip-flop driving a DSP48 input pin with *zero* logic levels and predominantly route delay, missing the target by 0.94 ns and costing 25 % of $F_{\max}$, while register count rises by 272. The cause is that DSP48E1 A/B/M pipeline registers support only synchronous reset, so an asynchronously reset register can never be absorbed into the slice. This constraint is documented [11] rather than new; what is worth reporting is that it is *decisive* for sparsity gating specifically. Gating an operand register is the only way to quiet a combinational multiplier, so a house reset style that merely costs area in an ordinary datapath here determines whether the entire technique is affordable.

Removing the reset from the datapath registers—which do not require one, since correctness is enforced by the validity and skip tags, and those *are* reset—restores full packing: $\text{AREG} = \text{BREG} = \text{MREG} = \text{PREG} = 8$ (final block). The accumulator's carry chain is absorbed into the DSP's own adder as well, halving CARRY4 usage from 127 to 63. The two variants are cycle-for-cycle bit-identical over 3031 cycles and 2823 result pulses, so this is a coding-style change with no architectural content—yet it decides whether the technique is affordable. The net outcome is that the gated, DSP-packed cascade is **33 % smaller and 4.9 % faster than the ungated baseline it replaces**, at the cost of one additional cycle of latency.

D. Power

Table III reports dynamic power at matched sparsity. Two properties of the measurement deserve stating, because both are routinely omitted.

First, *the ungated baseline's own power falls with sparsity*: zero operands toggle a multiplier array less regardless of any gating, and the baseline drops 31 % between $s = 0$ and $s = 0.9$ with no gating whatsoever. Comparing a gated design at high sparsity against an ungated one at $s = 0$ therefore credits the gating with a reduction that is mostly a property of the data. Each design here is compared against its own structural control at the same sparsity. On that basis split-enable zero-skip costs 3 % at $s = 0$ —the price of the detect and tag logic when there is no sparsity to exploit—and returns 26 % at $s = 0.75$ and 40 % at $s = 0.9$.

Second, *clock power is a floor*. No BUFGCE is inferred in any variant: enable-based gating on this device maps to the flip-flops' native CE pins, not to clock-tree gating. That

TABLE III
DYNAMIC POWER (mW) AT MATCHED OPERAND SPARSITY

Design	operand sparsity s (%)			
	0	50	75	90
Baseline (2-stage, async reset)	70	65	58	48
+ split-enable zero-skip	72	56	43	29
<i>change vs. its control</i>	+3%	-14%	-26%	-40%
DSP-packed, gating off	46	45	45	43
+ split-enable zero-skip	47	45	41	37
<i>change vs. its control</i>	+2%	-0%	-9%	-14%
clock component (baseline)	20	20	20	20

Vivado `report_power`; confidence Medium. Each design is compared against its own structural control at the same sparsity, because the ungated baseline itself loses 31 % of its power between $s = 0$ and $s = 0.9$ from reduced multiplier switching alone. The DSP-packed rows understate the gating benefit: see Section VI-E.

component is 20 mW in the baseline at every sparsity—29 % of its dynamic power—and no technique in this class removes it. A percentage saving quoted without stating this floor overstates what is achievable on FPGA.

Separately, DSP packing alone reduces dynamic power from 70 mW to 46 mW at $s = 0$, consistent with the area reduction of Table II. The reset-style correction is thus worth more in absolute terms than the gating itself, and the two are complementary rather than alternatives.

E. A limitation of SAIF-based comparison

The DSP-packed rows of Table III *understate* the benefit of gating, and we report them as a lower bound rather than a result. Applying an RTL-simulation SAIF to a routed checkpoint annotated only 12–13 % of design nets. For two designs of identical structure this is tolerable, because the shortfall is common-mode; between the fabric-register and DSP-packed designs it is not, because the gated operand registers of the latter are absorbed *into* the DSP48 slice and so have no corresponding net to annotate. The symptom is unmistakable: the DSP power component of the packed design reads a flat 20 mW at every sparsity while the baseline's falls from 21 mW to 7 mW. That flatness is an artefact of the annotation, not a property of the circuit. Post-implementation netlist simulation is the standard remedy and did not improve net matching in our attempts; resolving it is left as future work.

For that variant the two quantities independent of any power model are therefore the more informative: the implementation data of Table II, and the direct switching measurement of Section VI-B, in which multiplier-input transitions fall by 62 %.

F. Standard-cell corroboration

The FPGA numbers are entangled with a hard macro: the reset style decides whether the operand registers pack into the DSP48, which decides whether the multiplier can be quieted at all. To separate what is architectural from what is an artefact of that macro, all five designs were also synthesised to the open sky130_fd_sc_hd standard-cell library with Yosys and ABC (Table IV). There is no hard multiplier here; the array is ordinary gates.

TABLE IV
STANDARD-CELL SYNTHESIS, SKY130_FD_SC_HD (TYP., 25 °C, 1.8 V)

Design	area (μm^2)	vs. base	flip-flops	
			total	native-EN
Baseline (2-stage, async reset)	111,244	—	794	0
+ split-enable zero-skip	113,032	+1.6%	816	0
3-stage, gated operands, async reset	123,614	+11.1%	1088	0
3-stage, no datapath reset (control)	117,311	+5.5%	1058	512
+ split-enable zero-skip	119,088	+7.1%	1088	512

Yosys 0.63 + ABC, pre-layout cell area. “native-EN” counts registers mapped onto a true enable flip-flop (`edfextp`); the library provides an enable cell and an async-reset cell but none that is both, so asynchronously reset registers with a clock enable become a reset flop plus a feedback multiplexer instead.

Two things follow. First, *the gating is cheaper than the FPGA suggests*: it costs +1.6% of cell area on the two-stage pipeline and +1.5% against the three-stage control, against the +12.8% and +28% LUT figures of Table II. The added flip-flop count is identical on both targets (+22 and +30), so the discrepancy is LUT granularity inflating the apparent cost of small control cones, not a real difference.

Second, and less expected, *the reset-style penalty is not an FPGA artefact*. It reappears on the ASIC through an unrelated mechanism: the library offers an enable flip-flop and an asynchronous-reset flip-flop but no cell that is both, so an asynchronously reset register carrying a clock enable must be built as a reset flop plus a feedback multiplexer. The no-reset variants map 512 registers onto native enable cells; the asynchronously reset variants map none, and pay +3.8% area for the same function. A gating scheme is therefore only as good as the register style it is written in, on either target — which is the practical lesson we would most want a reader to take from this brief.

These are synthesis results only: no place and route, and no static timing, since no open STA tool was available in the environment used. Area is pre-layout cell area and the comparison is relative.

VII. CONCLUSION

Sparsity gating and completion-aware clock gating do not compose. The standard zero-skip enable applied to a drain-gated MAC pipeline destroys 63% of results, and the reason—a correctly sized drain window has no slack—is structural rather than incidental, so it will recur in any design combining the two. Splitting the enable into a never-masked control enable and a per-item-tagged datapath enable restores composability with no loss of throughput and no change to completion timing. Two secondary results are of independent interest. Register-enable activity counts substantially overstate the benefit of gating a multiplier’s output register, because gating a pipeline register downstream of a combinational block does not quiet that block. And the reset style in which a gated register is written decides whether the technique is affordable at all: asynchronous datapath resets cost 25% of F_{max} on the FPGA by blocking DSP48 absorption, and 3.8% of cell area on sky130 by forcing every enabled register into a reset flop with a feedback multiplexer. The mechanisms are unrelated;

the coding rule is the same. A gating scheme is only as good as the register style it is written in.

REFERENCES

- [1] L. Benini and G. De Micheli, “Automatic synthesis of low-power gated-clock finite-state machines,” *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 15, no. 6, pp. 630–643, Jun. 1996.
- [2] Q. Wu, M. Pedram, and X. Wu, “Clock-gating and its application to low power design of sequential circuits,” *IEEE Trans. Circuits Syst. I, Fundam. Theory Appl.*, vol. 47, no. 3, pp. 415–420, Mar. 2000.
- [3] Y.-H. Chen, T. Krishna, J. S. Emer, and V. Sze, “Eyeriss: An energy-efficient reconfigurable accelerator for deep convolutional neural networks,” *IEEE J. Solid-State Circuits*, vol. 52, no. 1, pp. 127–138, Jan. 2017.
- [4] J. Albericio, P. Judd, T. Hetherington, T. Aamodt, N. E. Jerger, and A. Moshovos, “Cnvlutin: Ineffectual-neuron-free deep neural network computing,” in *Proc. ACM/IEEE Int. Symp. Comput. Archit. (ISCA)*, 2016, pp. 1–13.
- [5] A. Parashar *et al.*, “SCNN: An accelerator for compressed-sparse convolutional neural networks,” in *Proc. ACM/IEEE Int. Symp. Comput. Archit. (ISCA)*, 2017, pp. 27–40.
- [6] S. Han *et al.*, “EIE: Efficient inference engine on compressed deep neural network,” in *Proc. ACM/IEEE Int. Symp. Comput. Archit. (ISCA)*, 2016, pp. 243–254.
- [7] D. Gohel, A. Gundrapally, and K. Choi, “RTL-level power optimization of CNN accelerators via clock gating and sparsity-aware MAC suppression on FPGA,” *Electronics*, vol. 15, no. 11, art. 2492, 2026.
- [8] O. Temam *et al.*, “Exploiting input data sparsity in neural network compute units,” U.S. Patent 9818 059, Nov. 14, 2017.
- [9] Imagination Technologies Ltd., “Hardware unit for performing matrix multiplication with clock gating,” Chinese Patent CN110 007 896B, May 13, 2022.
- [10] International Business Machines Corp., “Interlocked synchronous pipeline clock gating,” U.S. Patent 7 308 593, Dec. 11, 2007.
- [11] *Vivado Design Suite User Guide: Synthesis*, UG901, Xilinx, Inc., San Jose, CA, USA, 2023.
- [12] J. Park *et al.*, “Sparsity-aware and re-configurable NPU architecture for Samsung flagship mobile SoC,” in *Proc. ACM/IEEE Int. Symp. Comput. Archit. (ISCA)*, 2021, pp. 15–28.